

System Programming

Lecture 1

Yeongpil Cho

Hanyang University

About me

■ 조영필 (Yeongpil Cho)

- A system security researcher
- Designing new SW/HW techniques for better security
 - Kernels & Firmware
 - Applications
 - (co-)Processors

■ Contact

- ypcho@hanyang.ac.kr
- TA: 강범우
 - qjadrn0630@hanyang.ac.kr



Objectives

- **Computer systems consists of various hardware and software components.**
- **Most system implementations share some fundamental design principles.**
 - e.g., Android Phone, iPhone, Windows desktop
- **We are going to generally explore those principles**
 - to understand how programs run on the system
 - to understand what affect programs
 - to, ultimately, develop better programs

Materials

■ Textbook

- Computer Systems: A Programmer's Perspective 3rd Edition
 - Randal E Bryant, David R O'Hallaron

■ Reference Course

- Carnegie Mellon University's Intro to Computer Systems

Evaluation

- **Exam** **60%**
 - Midterm 30%
 - Final 30%
- **Lab** **30%**
- **Attendance** **10%**
 - 3 tardies → 1 absence

Topics

- Overview of the system programming course

Abstraction Is Good But Don't Forget Reality

- **Most CS and CE courses emphasize abstraction**
 - Abstract Data Type
 - List, Tree, ...
 - Asymptotic analysis
 - Big-O/ Θ / Ω , ...
- **These abstractions are good for understanding, but we need to know reality**
 - for better performance optimization
 - for better secure implementation

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

■ Float's: Yes!


■ Int's:

■ $40,000 * 40,000 \rightarrow 1,600,000,000$

■ $50,000 * 50,000 \rightarrow 2,500,000,000 ?? \rightarrow \text{No!} \rightarrow -1,794,967,296$

– Range of Integer value

» $-2,147,483,648 \sim 2,147,483,647$



Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 2: Is $(x + y) + z = x + (y + z)$?

- Unsigned & Signed Int's: Yes!

- Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$

- $1e20 + (-1e20 + 3.14) \rightarrow 3.14 ?? \rightarrow \text{No!} \rightarrow 0$

- Why?

- Addition of float numbers

- » (1) equalize exponents to the larger one

- » (2) adjust mantissa according to the modified exponent

- $1e20 \rightarrow 0b0\ 11000001\ 01011010111100011101100$

- $-1e20 \rightarrow 0b1\ 11000001\ 01011010111100011101100$

- $3.14 \rightarrow 0b0\ 10000000\ 10010001111010111000011$

Computer Arithmetic

■ Does not generate random values

- Arithmetic operations have important mathematical properties

■ Cannot assume all “usual” mathematical properties

- Due to finiteness of representations
- Integer operations satisfy “ring” properties
 - Commutativity, associativity, distributivity
- Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs

■ Observation

- Need to understand which abstractions apply in which contexts
- Important issues for compiler writers and serious application programmers

Great Reality #2:

You've Got to Know Assembly

- **Chances are, you'll never write programs in assembly**
 - Compilers are much better & more patient than you are
- **But: Understanding assembly is key to machine-level execution model**
 - Behavior of programs in presence of bugs
 - High-level language models break down
 - Tuning program performance
 - Understand optimizations done / not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting malware
 - x86 assembly is the language of choice!

Assembly Code Example

■ Time Stamp Counter

- a 64-bit register on x86
- increment on every CPU clock
- readable using `rdtsc` instruction

■ Elapsed time of a procedure.

```
unsigned long long t1, t2;  
t1 = read_counter();  
P();  
t2 = read_counter();  
  
printf("execution time of P():%d clock cycles\n", t2-t1);
```

Assembly Code Example

- Implementing through inline assembly code
- Or, linking an object file of an assembly code

```
unsigned long long read_counter() {  
    unsigned long long hi, lo;  
  
    asm("rdtsc          \n  
        movl %%edx, %0  \n  
        movl %%eax, %1  \n")  
        : "=r" (hi), "=r" lo : : "%edx", "%eax");  
        return (lo | (hi << 32));  
}
```

Great Reality #3: Memory Matters

Random Access Memory Is an Unphysical Abstraction

■ Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated

■ Memory referencing bugs especially pernicious

- Effects are distant in both time and space

■ Memory performance is not uniform

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

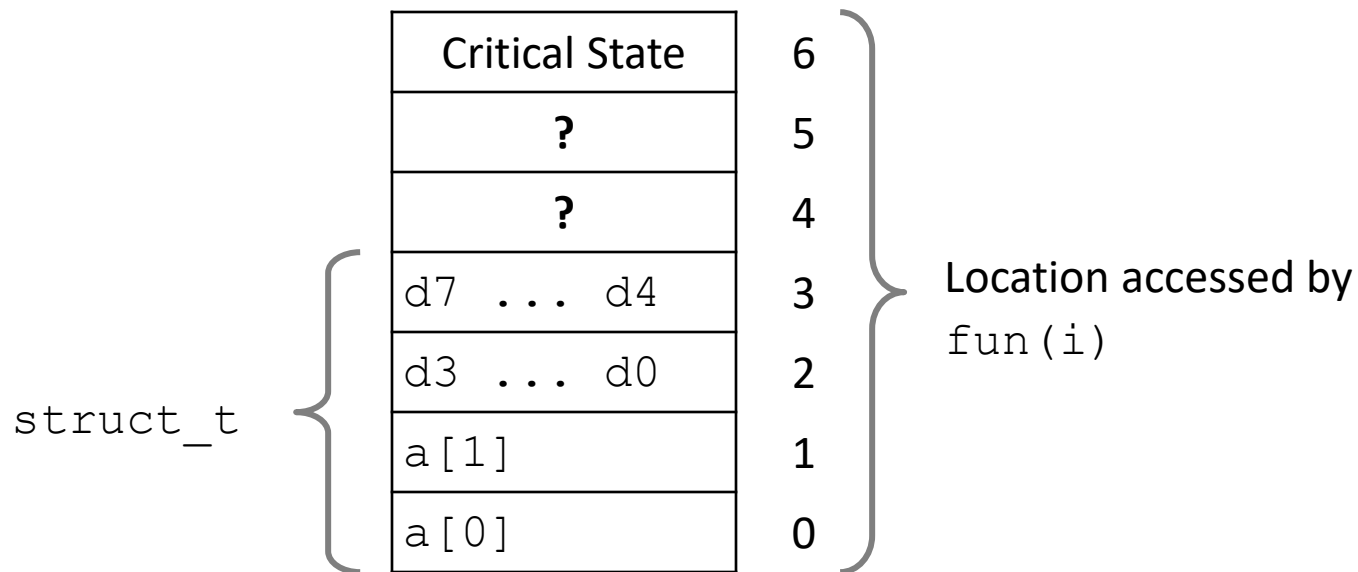
fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.13999998664856
fun(3)	→	2.000000061035156
fun(4)	→	3.14
fun(6)	→	Segmentation fault

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

fun(0) → 3.14
fun(1) → 3.14
fun(2) → 3.1399998664856
fun(3) → 2.00000061035156
fun(4) → 3.14
fun(6) → Segmentation fault

Explanation:



Memory Referencing Errors

■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

■ Can lead to nasty bugs

- Whether or not bug has any effect depends on system and compiler
- Action at a distance
 - Corrupted object logically unrelated to one being accessed
 - Effect of bug may be first observed long after it is generated

■ How can I deal with this?

- Program in Java, Ruby, Python, ML, ...
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. Valgrind)

Great Reality #4: There's more to performance than asymptotic complexity

- **Constant factors matter too!**
- **And even exact op count does not predict performance**
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- **Must understand system to optimize performance**
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

4.3ms


2.0 GHz Intel Core i7 Haswell

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

81.8ms

- Hierarchical memory organization
- Performance depends on access patterns
 - Including how step through multi-dimensional array

Why The Performance Differs



Great Reality #5:

Computers do more than execute programs

- **They need to get data in and out**
 - I/O system critical to program reliability and performance
- **They communicate with each other over networks**
 - Many system-level issues arise in presence of network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues

Amdahl's Law

■ When we improve a part of the system, how does it affect the entire performance?

- T_{old} = Original execution time
- T_{new} = New execution time
- α = Ratio of the improved part
- k = Degree of improvement

- $T_{new} = (1 - \alpha)T_{old} + (\alpha T_{old})/k$

- (Speed up) $S = T_{old}/T_{new} = \frac{1}{(1-\alpha)+\alpha/k}$

- ex1) improve 60% of the system ($\alpha = 0.6$) by three times ($k = 3$)

- $S = 1/[(1-0.6) + 0.6/3] = 1.67$

- ex2) improve 30% of the system ($\alpha = 0.3$) by twenty times ($k = 20$)

- $S = 1/[(1-0.3) + 0.3/20] = 1.4$

■ Conclusion: Try to improve a major part of the system as possible.